

```
In [2]: # To access Google Drive:
from google.colab import drive

drive.mount("/content/gdrive")
```

Mounted at /content/gdrive

```
In [3]: REVIEW_PATH = "/content/gdrive/MyDrive/reviews.csv"
```

```
In [4]: import pandas as pd
import numpy as np
import tensorflow as tf
```

```
In [5]: BATCH_SIZE=8
```

```
In [6]: data = pd.read_csv(REVIEW_PATH)
```

```
In [7]: data.head()
```

```
Out[7]:
```

	reviewId	userName
0	gp:AOqpTOEhZuqSqqWnaKRgv-9ABYdajFUB0WugPGh-SG-...	Eric Tie lh.googleus
1	gp:AOqpTOH0WP4IQKBZ2LrdNmFy_YmpPCVrV3diEU9KGm3...	john alpha lh.googleus
2	gp:AOqpTOEMCkJB8lq1p-r9dPwnSYadA5BkPWTf32Z1azu...	Sudhakar .S lh.googleus
3	gp:AOqpTOGFrUWuKGycpje8kszj3uwHN6tU_fd4gLVFy9z...	SKGflorida@bellsouth.net DAVID S lh.googleus
4	gp:AOqpTOHls7DW8wmDFzTkHwxuqFkdNQtkHmO6Pt9jhZE...	Louann Stoker https://play-l

```
In [7]:
```

```
In [8]: data = data[["content", "score"]].rename(columns={"content": "text", "score": "label"})
```

```
In [9]: data.head()
```

```
Out[9]:
```

	text	label
0	I cannot open the app anymore	1
1	I have been begging for a refund from this app...	1
2	Very costly for the premium version (approx In...	1
3	Used to keep me organized, but all the 2020 UP...	1
4	Dan Birthday Oct 28	1

```
In [10]: data['label'] = data['label'].replace([1,2,3,4,5],[0,1,2,3,4])
```

```
In [11]: data['label'].value_counts()
```

```
Out[11]: label
4      2879
3      2775
0      2506
1      2344
2      1991
Name: count, dtype: int64
```

```
In [12]: !pip install datasets
```

## Collecting datasets

```
Downloading datasets-2.18.0-py3-none-any.whl (510 kB)
_____ 510.5/510.5 kB 3.9 MB/s eta 0:0
0:00
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-pac
kages (from datasets) (3.13.4)
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.10/dist-
packages (from datasets) (1.25.2)
Requirement already satisfied: pyarrow>=12.0.0 in /usr/local/lib/python3.10/di
st-packages (from datasets) (14.0.2)
Requirement already satisfied: pyarrow-hotfix in /usr/local/lib/python3.10/di
st-packages (from datasets) (0.6)
Collecting dill<0.3.9,>=0.3.0 (from datasets)
  Downloading dill-0.3.8-py3-none-any.whl (116 kB)
  _____ 116.3/116.3 kB 15.4 MB/s eta 0:
00:00
Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packa
ges (from datasets) (2.0.3)
Requirement already satisfied: requests>=2.19.0 in /usr/local/lib/python3.10/
dist-packages (from datasets) (2.31.0)
Requirement already satisfied: tqdm>=4.62.1 in /usr/local/lib/python3.10/dist
-packages (from datasets) (4.66.2)
Collecting xxhash (from datasets)
  Downloading xxhash-3.4.1-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x8
6_64.whl (194 kB)
  _____ 194.1/194.1 kB 23.8 MB/s eta 0:
00:00
Collecting multiprocessing (from datasets)
  Downloading multiprocessing-0.70.16-py310-none-any.whl (134 kB)
  _____ 134.8/134.8 kB 17.7 MB/s eta 0:
00:00
Requirement already satisfied: fsspec[http]<=2024.2.0,>=2023.1.0 in /usr/loca
l/lib/python3.10/dist-packages (from datasets) (2023.6.0)
Requirement already satisfied: aiohttp in /usr/local/lib/python3.10/dist-pack
ages (from datasets) (3.9.3)
Requirement already satisfied: huggingface-hub>=0.19.4 in /usr/local/lib/pyth
on3.10/dist-packages (from datasets) (0.20.3)
Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-pa
ckages (from datasets) (24.0)
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.10/dist-
packages (from datasets) (6.0.1)
Requirement already satisfied: aiosignal>=1.1.2 in /usr/local/lib/python3.10/
dist-packages (from aiohttp->datasets) (1.3.1)
Requirement already satisfied: attrs>=17.3.0 in /usr/local/lib/python3.10/dis
t-packages (from aiohttp->datasets) (23.2.0)
Requirement already satisfied: frozenlist>=1.1.1 in /usr/local/lib/python3.1
0/dist-packages (from aiohttp->datasets) (1.4.1)
Requirement already satisfied: multidict<7.0,>=4.5 in /usr/local/lib/python3.
10/dist-packages (from aiohttp->datasets) (6.0.5)
Requirement already satisfied: yarll<2.0,>=1.0 in /usr/local/lib/python3.10/di
st-packages (from aiohttp->datasets) (1.9.4)
Requirement already satisfied: async-timeout<5.0,>=4.0 in /usr/local/lib/pyth
on3.10/dist-packages (from aiohttp->datasets) (4.0.3)
Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/p
ython3.10/dist-packages (from huggingface-hub>=0.19.4->datasets) (4.11.0)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/pyt
hon3.10/dist-packages (from requests>=2.19.0->datasets) (3.3.2)
```

Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests>=2.19.0->datasets) (3.6)  
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests>=2.19.0->datasets) (2.0.7)  
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests>=2.19.0->datasets) (2024.2.2)  
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.10/dist-packages (from pandas->datasets) (2.8.2)  
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas->datasets) (2023.4)  
Requirement already satisfied: tzdata>=2022.1 in /usr/local/lib/python3.10/dist-packages (from pandas->datasets) (2024.1)  
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.8.2->pandas->datasets) (1.16.0)  
Installing collected packages: xxhash, dill, multiprocessing, datasets  
Successfully installed datasets-2.18.0 dill-0.3.8 multiprocessing-0.70.16 xxhash-3.4.1

In [12]:

## Train Test Split

In [13]: `from sklearn.model_selection import train_test_split`

```
train_data, test_data = train_test_split(data, test_size=0.2)
```

In [14]: `from datasets import DatasetDict, Dataset`

In [15]: `train_dataset = Dataset.from_pandas(train_data)`  
`test_dataset = Dataset.from_pandas(test_data)`  
  
*# Create a DatasetDict to store both train and test datasets*  
`dataset = DatasetDict({'train': train_dataset, 'test': test_dataset})`

In [16]: `dataset`

Out[16]: `DatasetDict({`  
    `train: Dataset({`  
        `features: ['text', 'label', '__index_level_0__'],`  
        `num_rows: 9996`  
    `})`  
    `test: Dataset({`  
        `features: ['text', 'label', '__index_level_0__'],`  
        `num_rows: 2499`  
    `})`  
`})`

# Loading Tokenizer

```
In [17]: from transformers import (BertTokenizerFast,TFBertTokenizer,BertTokenizer,Robe
        DataCollatorWithPadding,TFRobertaForSequenceClassifi
        TFBertModel,create_optimizer)
from transformers import TFAutoModel, AutoTokenizer
```

```
In [18]: # model_id="nlptown/bert-base-multilingual-uncased-sentiment"
# tokenizer = BertTokenizerFast.from_pretrained(model_id)
tokenizer = AutoTokenizer.from_pretrained("bert-base-uncased")
```

/usr/local/lib/python3.10/dist-packages/huggingface\_hub/utils/\_token.py:88: UserWarning:  
The secret `HF\_TOKEN` does not exist in your Colab secrets.  
To authenticate with the Hugging Face Hub, create a token in your settings tab (https://huggingface.co/settings/tokens), set it as secret in your Google Colab and restart your session.  
You will be able to reuse this secret in all of your notebooks.  
Please note that authentication is recommended but still optional to access public models or datasets.

```
warnings.warn(

tokenizer_config.json:  0%|          | 0.00/48.0 [00:00<?, ?B/s]
config.json:          0%|          | 0.00/570 [00:00<?, ?B/s]
vocab.txt:            0%|          | 0.00/232k [00:00<?, ?B/s]
tokenizer.json:       0%|          | 0.00/466k [00:00<?, ?B/s]
```

## Passing review data into tokenizer

```
In [19]: output=tokenizer(dataset['train'][:2]['text'],padding=True,truncation=True,max
# # print(output)
```

## Checking the decoded version to tokenized data

Here the [CLS] shows the start of the sentence and [SEP] shows the end of the sentence [PAD] is added to end the sequence to make it reach the maximum length. in other words it ensure that all of the sequence are of the same length

```
In [20]: tokenizer.decode(output['input_ids'][0])
```

```
Out[20]: '[CLS] it was good until they updated it and took out the today, someday, upc
[PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD]
[PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD]
[PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD]
[PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD]
[PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD]
[PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD]
[PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD]'
```

### Applying to whole dataset

```
In [21]: def preprocess_function(examples):
         return tokenizer(examples["text"],padding=True,truncation=True,max_length=12
```

```
In [22]: tokenized_dataset = dataset.map(preprocess_function, batched=True)
         tokenized_dataset
```

```
Map:   0%|          | 0/9996 [00:00<?, ? examples/s]
```

```
Map:   0%|          | 0/2499 [00:00<?, ? examples/s]
```

```
Out[22]: DatasetDict({
  train: Dataset({
    features: ['text', 'label', '__index_level_0__', 'input_ids', 'token_
type_ids', 'attention_mask'],
    num_rows: 9996
  })
  test: Dataset({
    features: ['text', 'label', '__index_level_0__', 'input_ids', 'token_
type_ids', 'attention_mask'],
    num_rows: 2499
  })
})
```

### Creating tensorflow dataset from tokenized dataset

```
In [23]: tf_train_datasets = tokenized_dataset["train"].to_tf_dataset(
         columns=['input_ids', 'token_type_ids', 'attention_mask', 'label'],
         shuffle=True,
         batch_size=BATCH_SIZE,
         )
```

```
In [24]: tf_val_datasets = tokenized_dataset["test"].to_tf_dataset(  
    columns=['input_ids', 'token_type_ids', 'attention_mask', 'label'],  
    shuffle=True,  
    batch_size=BATCH_SIZE,  
    #collate_fn=data_collator  
)
```

Separating the Labels from the dataset

```
In [25]: def swap_positions(dataset):  
    return {'input_ids':dataset['input_ids'],  
        'token_type_ids':dataset['token_type_ids'],  
        'attention_mask':dataset['attention_mask'],},dataset['label']
```

```
In [26]: tf_train_dataset=tf_train_datasets.map(swap_positions).prefetch(tf.data.AUTOTUNE)  
tf_val_dataset=tf_val_datasets.map(swap_positions).prefetch(tf.data.AUTOTUNE)
```

## Modelling Using TFBertForSequenceClassification

```
In [27]: # from transformers import AutoModelForSequenceClassification
```



```
In [28]: # model=TFBertForSequenceClassification.from_pretrained("nlpTown/bert-base-multisentenceclassification")
model = TFAutoModel.from_pretrained("bert-base-uncased")
model.summary()
```

```
model.safetensors: 0%|          | 0.00/440M [00:00<?, ?B/s]
```

Some weights of the PyTorch model were not used when initializing the TF 2.0 model TFBertModel: ['cls.predictions.bias', 'cls.seq\_relationship.weight', 'cls.predictions.transform.dense.weight', 'cls.predictions.transform.LayerNorm.bias', 'cls.seq\_relationship.bias', 'cls.predictions.transform.dense.bias', 'cls.predictions.transform.LayerNorm.weight']

- This IS expected if you are initializing TFBertModel from a PyTorch model trained on another task or with another architecture (e.g. initializing a TFBertForSequenceClassification model from a BertForPreTraining model).

- This IS NOT expected if you are initializing TFBertModel from a PyTorch model that you expect to be exactly identical (e.g. initializing a TFBertForSequenceClassification model from a BertForSequenceClassification model).

All the weights of TFBertModel were initialized from the PyTorch model.

If your task is similar to the task the model of the checkpoint was trained on, you can already use TFBertModel for predictions without further training.

Model: "tf\_bert\_model"

Layer (type)	Output Shape	Param #
bert (TFBertMainLayer)	multiple	109482240

```
=====  
Total params: 109482240 (417.64 MB)  
Trainable params: 109482240 (417.64 MB)  
Non-trainable params: 0 (0.00 Byte)
```

```
In [29]: class BERTForClassification(tf.keras.Model):

    def __init__(self, bert_model, num_classes):
        super().__init__()
        self.bert = bert_model
        # self.drop = tf.keras.layers.Dropout(rate=0.3)
        self.fc = tf.keras.layers.Dense(num_classes, activation='softmax')

    def call(self, inputs):
        x = self.bert(inputs)[1]
        # pooled_output = outputs[1] # Using [CLS] token output
        # output = self.dropout(pooled_output)
        return self.fc(x)
```

## Training

```
In [30]: # num_epochs = 5
# batches_per_epoch = len(tokenized_dataset["train"]) // BATCH_SIZE
# total_train_steps = int(batches_per_epoch * num_epochs)
```

```
In [31]: # optimizer, schedule = create_optimizer(init_lr=2e-5,num_warmup_steps=0, num_
```



```
In [32]: # print(optimizer)
```

<tf\_keras.src.optimizers.adam.Adam object at 0x7d3a904fb0d0>

```
In [33]: classifier = BERTForClassification(model, num_classes=5)

classifier.compile(
    optimizer=tf.keras.optimizers.Adam(),
    loss=tf.keras.losses.SparseCategoricalCrossentropy(),
    metrics=['accuracy']
)
```

```
In [ ]: history=classifier.fit(
    tf_train_dataset,
    validation_data=tf_val_dataset,
    epochs=3)
```

Epoch 1/3  
232/1250 [====>.....] - ETA: 3:15:08 - loss: 1.7368 - ac  
curacy: 0.2064

```
In [ ]: # model.compile(loss= tf.keras.Losses.SparseCategoricalCrossentropy(),
# optimizer=optimizer,
#     metrics=['accuracy'],)
#     #run_eagerly=True)

#     # CategoricalCrossentropy
```

```
In [ ]: # history=model.fit(
#     tf_train_dataset,
#     validation_data=tf_val_dataset,
#     epochs=5)
```

```
In [ ]: import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix, roc_curve
import seaborn as sns
```

```
In [ ]: plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model_loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'val'], loc='upper left')
plt.show()
```

```
In [ ]: plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])

plt.title('model_accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'val'], loc='upper left')
plt.show()
```

```
In [ ]: classifier.evaluate(tf_val_dataset)
```

```
In [ ]:
```